



Sorting in Linear Time





Comparison Sorting Review

- Insertion sort:
 - Pro's:
 - Easy to code
 - Fast on small inputs (less than ~50 elements)
 - Fast on nearly-sorted inputs
 - Con's:
 - $O(n^2)$ worst case
 - $O(n^2)$ average case
 - $O(n^2)$ reverse-sorted case



Comparison Sorting Review

- Merge sort:
 - Divide-and-conquer:
 - Split array in half
 - Recursively sort sub-arrays
 - Linear-time merge step
 - Pro's:
 - $O(n \lg n)$ worst case - asymptotically optimal for comparison sorts
 - Con's:
 - Doesn't sort in place



Comparison Sorting Review

- Heap sort:
 - Uses the very useful heap data structure
 - Complete binary tree
 - Heap property: parent key $>$ children's keys
 - Pro's:
 - $O(n \lg n)$ worst case - asymptotically optimal for comparison sorts
 - Sorts in place
 - Con's:
 - Fair amount of shuffling memory around



Comparison Sorting Review

- Quick sort:
 - Divide-and-conquer:
 - Partition array into two sub-arrays, recursively sort
 - All of first sub-array < all of second sub-array
 - Pro's:
 - $O(n \lg n)$ average case
 - Sorts in place
 - Fast in practice (why?)
 - Con's:
 - $O(n^2)$ worst case
 - Naïve implementation: worst case on sorted input
 - Good partitioning makes this very unlikely.



Non-Comparison Based Sorting

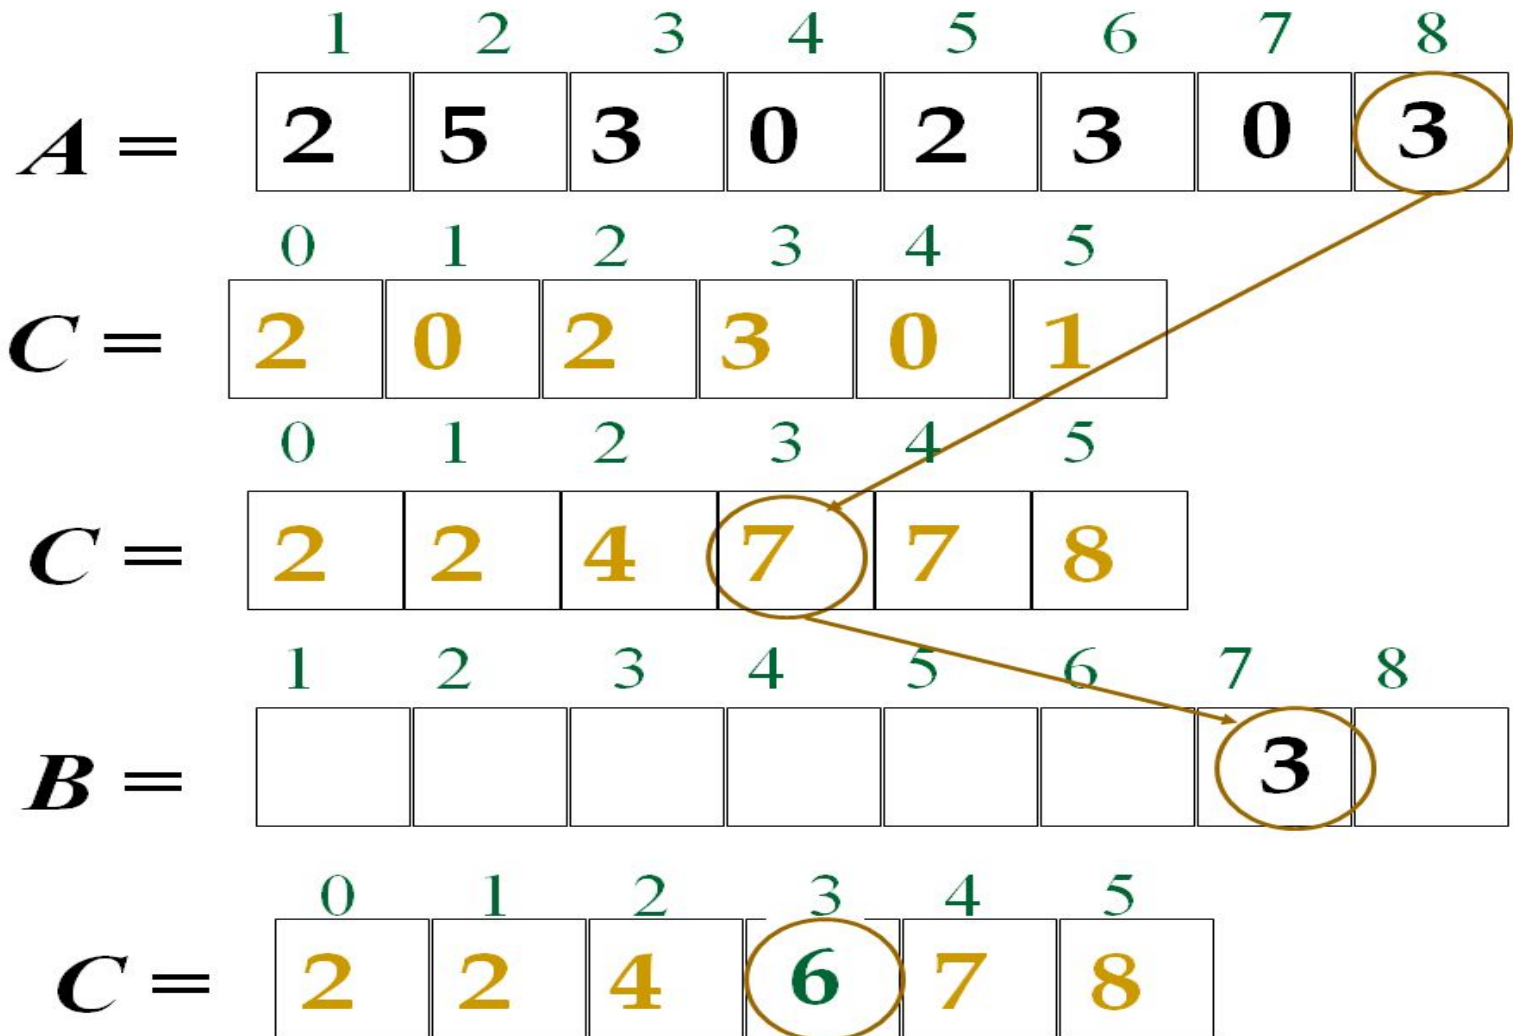
- Many times we have restrictions on our keys
 - Social Security Numbers
 - Employee ID's
- We will examine three algorithms which under certain conditions can run in $O(n)$ time.
 - Counting sort
 - Radix sort
 - Bucket sort



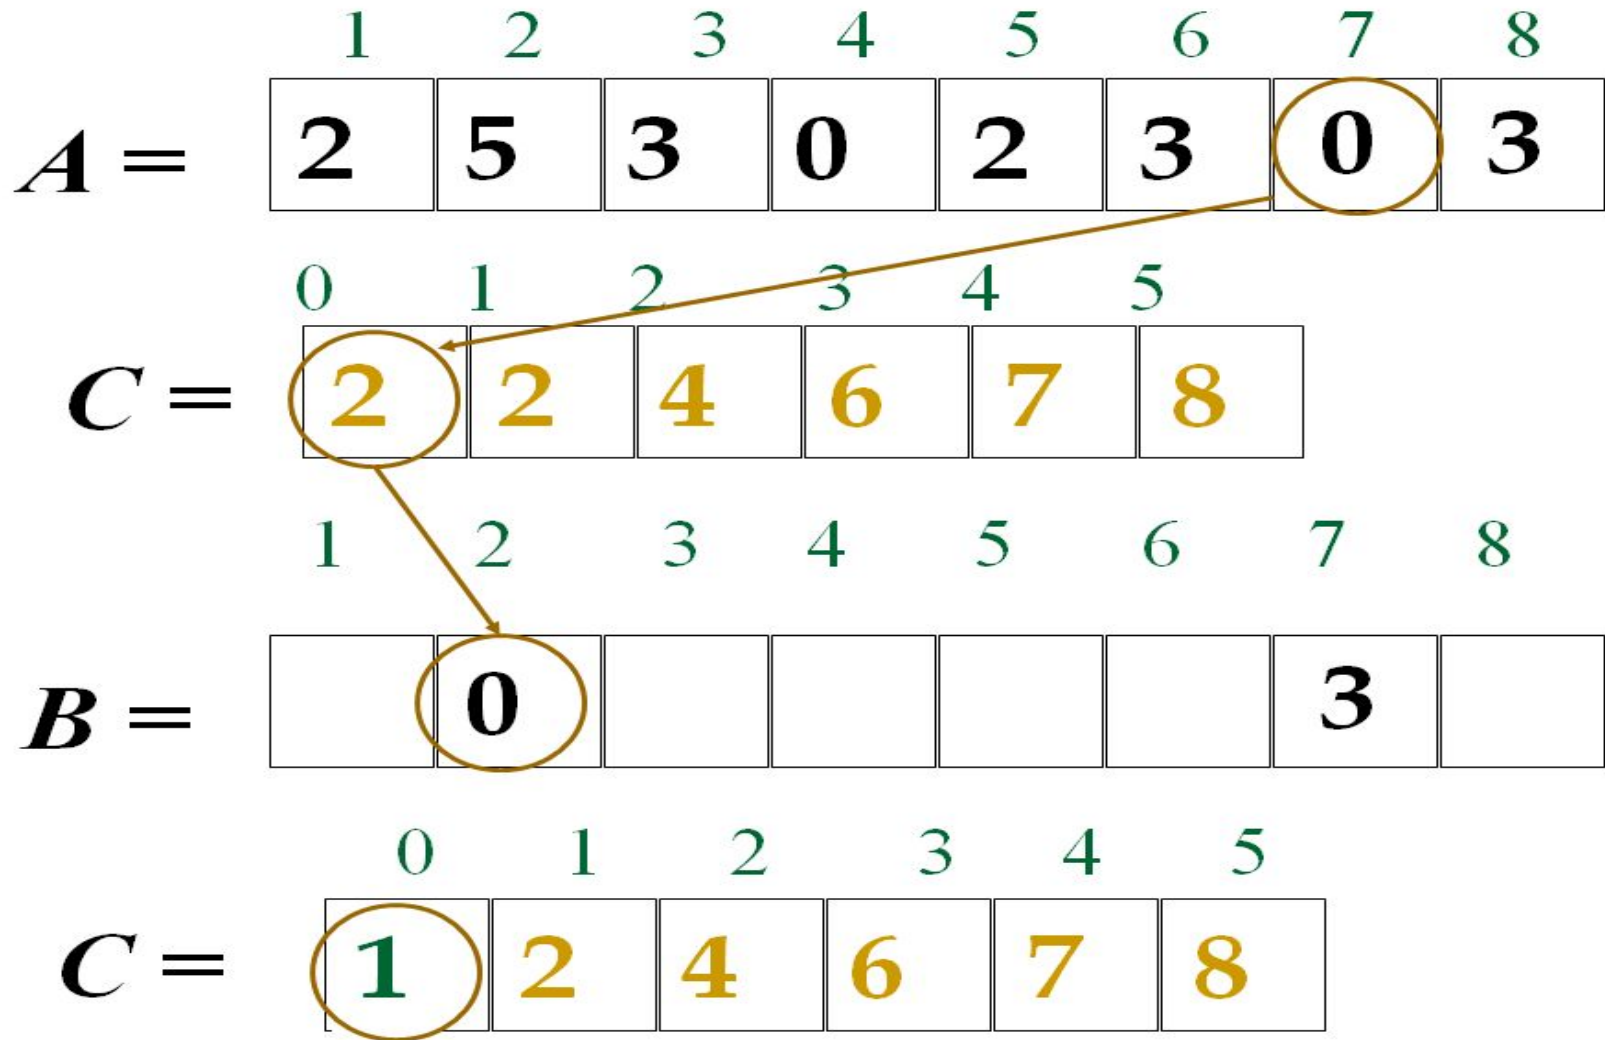
Counting Sort

- Why it's not a comparison sort:
 - Assumption: input - integers in the range $0..k$
 - No comparisons made!
- Basic idea:
 - determine for each input element x its rank: the number of elements less than x .
 - once we know the rank r of x , we can place it in position $r+1$
- Depends on assumption about the numbers being sorted
 - Assume numbers are in the range $1..k$
- The algorithm:
 - Input: $A[1..n]$, where $A[j] \in \{1, 2, 3, \dots, k\}$
 - Output: $B[1..n]$, sorted (not sorted in place)
 - Also: Array $C[1..k]$ for auxiliary storage

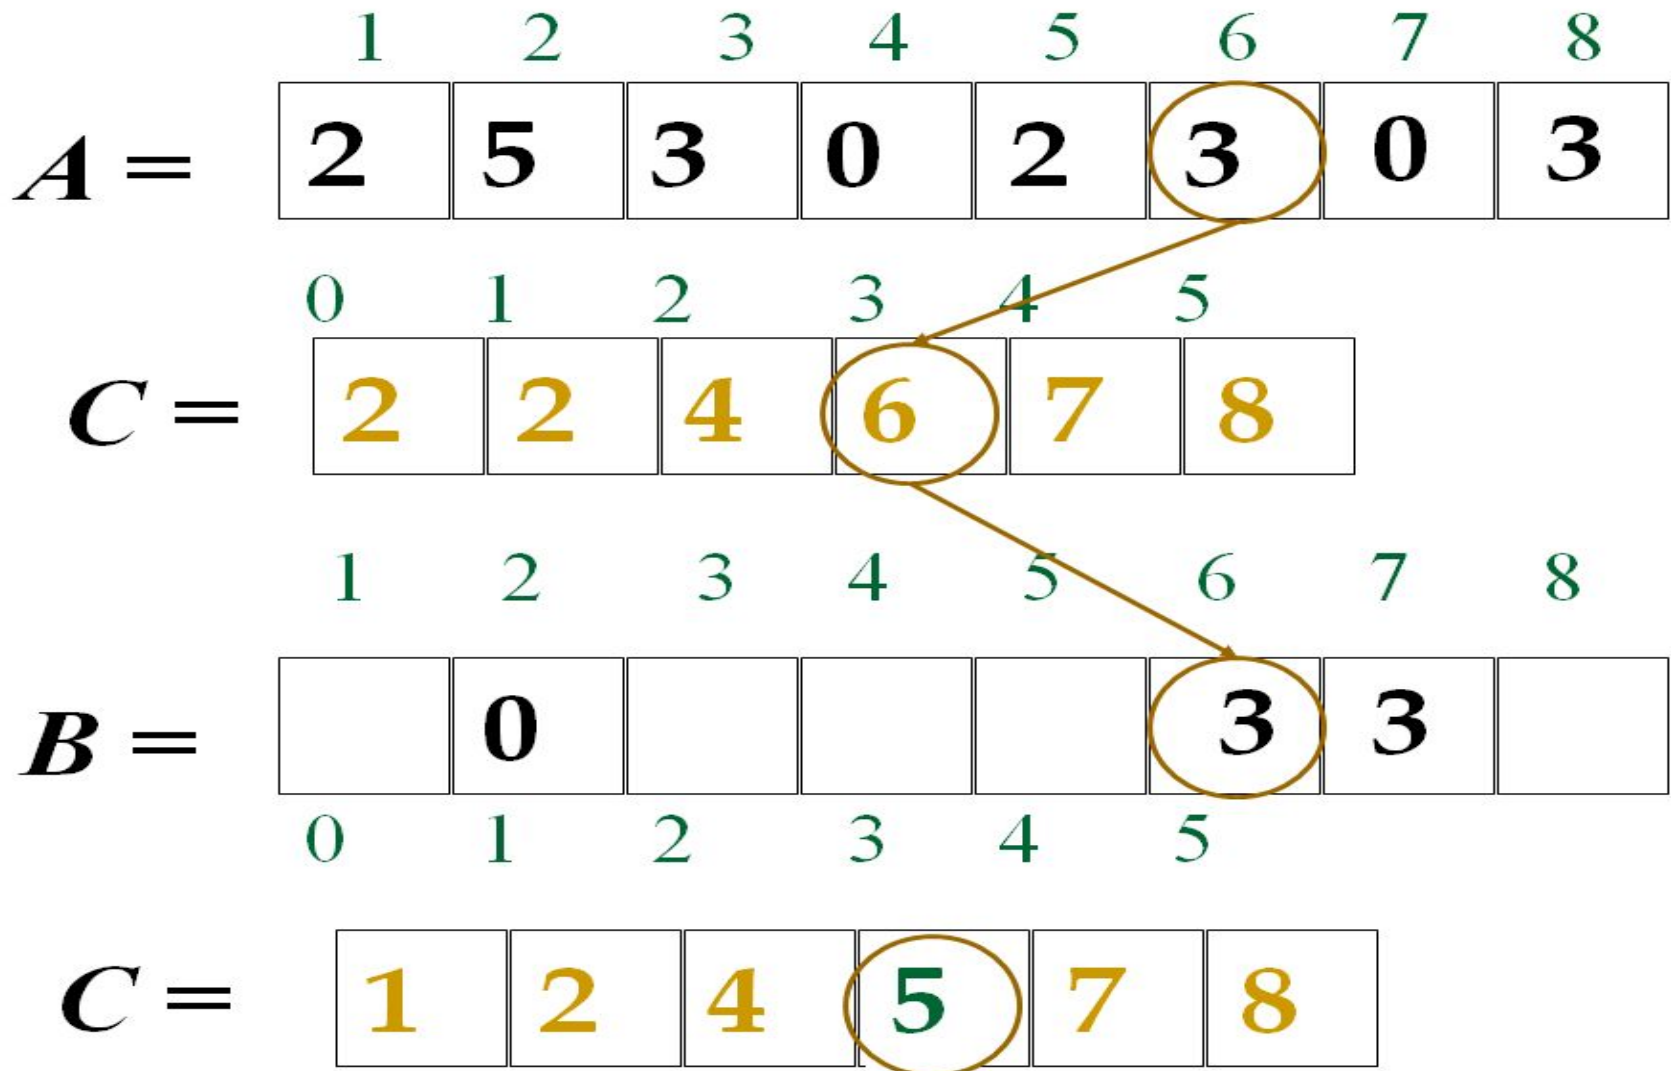
Counting Sort Example



Counting Sort Example



Counting Sort Example



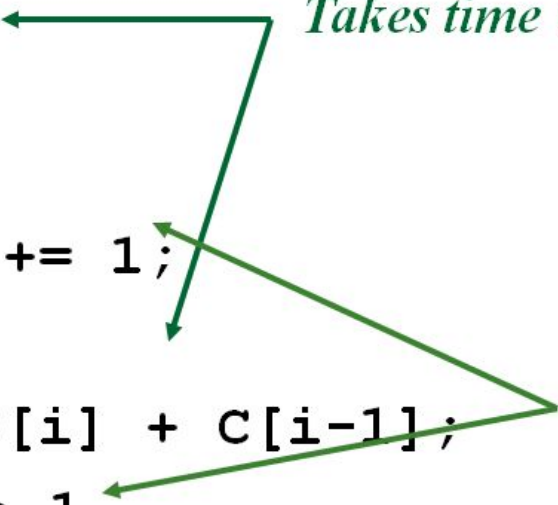


Counting Sort

```
1  CountingSort(A, B, k)
2      for i=1 to k
3          C[i] = 0;
4      for j=1 to n
5          C[A[j]] += 1;
6      for i=2 to k
7          C[i] = C[i] + C[i-1];
8      for j=n downto 1
9          B[C[A[j]]] = A[j];
10     C[A[j]] -= 1;
```

Takes time $O(k)$

Takes time $O(n)$



What will be the running time?



Counting Sort

- Total time: $O(n + k)$
 - Usually, $k = O(n)$
 - Thus counting sort runs in $O(n)$ time
- But sorting is $(n \lg n)!$
 - No contradiction--this is not a comparison sort (in fact, there are no comparisons at all!)
 - Notice that this algorithm is stable
 - If numbers have the same value, they keep their original order



Stable Sorting Algorithms

- A sorting algorithm is stable if for any two indices i and j with $i < j$ and $a_i = a_j$, element a_i precedes element a_j in the output sequence.

Input

2 ₁	7 ₁	4 ₁	4 ₂	2 ₂	5 ₁	2 ₃	6 ₁
----------------	----------------	----------------	----------------	----------------	----------------	----------------	----------------

Output

2 ₁	2 ₂	2 ₃	4 ₁	4 ₂	5 ₁	6 ₁	7 ₁
----------------	----------------	----------------	----------------	----------------	----------------	----------------	----------------

Observation: *Counting Sort is stable.*



Counting Sort

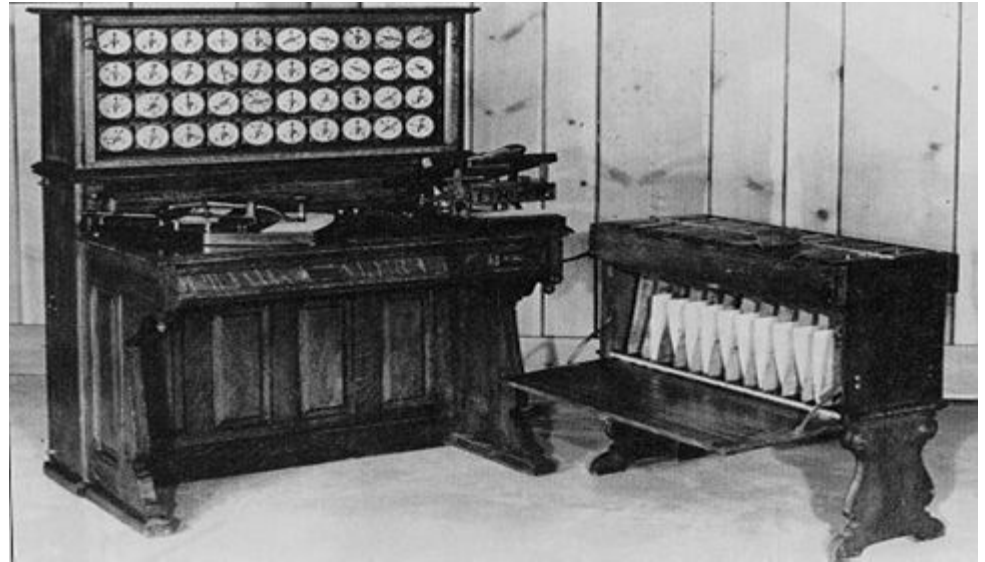
- ❑ Linear Sort! Cool! Why don't we always use counting sort?
- ❑ Because it depends on range k of elements
- ❑ Could we use counting sort to sort 32 bit integers? Why or why not?
- ❑ Answer: no, k too large ($2^{32} = 4,294,967,296$)

Counting Sort

- Assumption: input taken from small set of numbers of size k
- Basic idea:
 - Count number of elements less than you for each element.
 - This gives the position of that number – similar to selection sort.
- Pro's:
 - Fast
 - Asymptotically fast - $O(n+k)$
 - Simple to code
- Con's:
 - Doesn't sort in place.
 - Elements must be **integers**. *countable*
 - Requires $O(n+k)$ extra storage.

Radix Sort

- ❑ Origin : Herman Hollerith's card-sorting machine for the 1890 U.S Census



- ❑ Digit-by-digit sort
- ❑ Hollerith's original (bad) idea : sort on most-significant digit first.
- ❑ Good idea : Sort on least-significant digit first with auxiliary stable sort

Radix Sort

IBM 083
punch card
sorter





Radix Sort

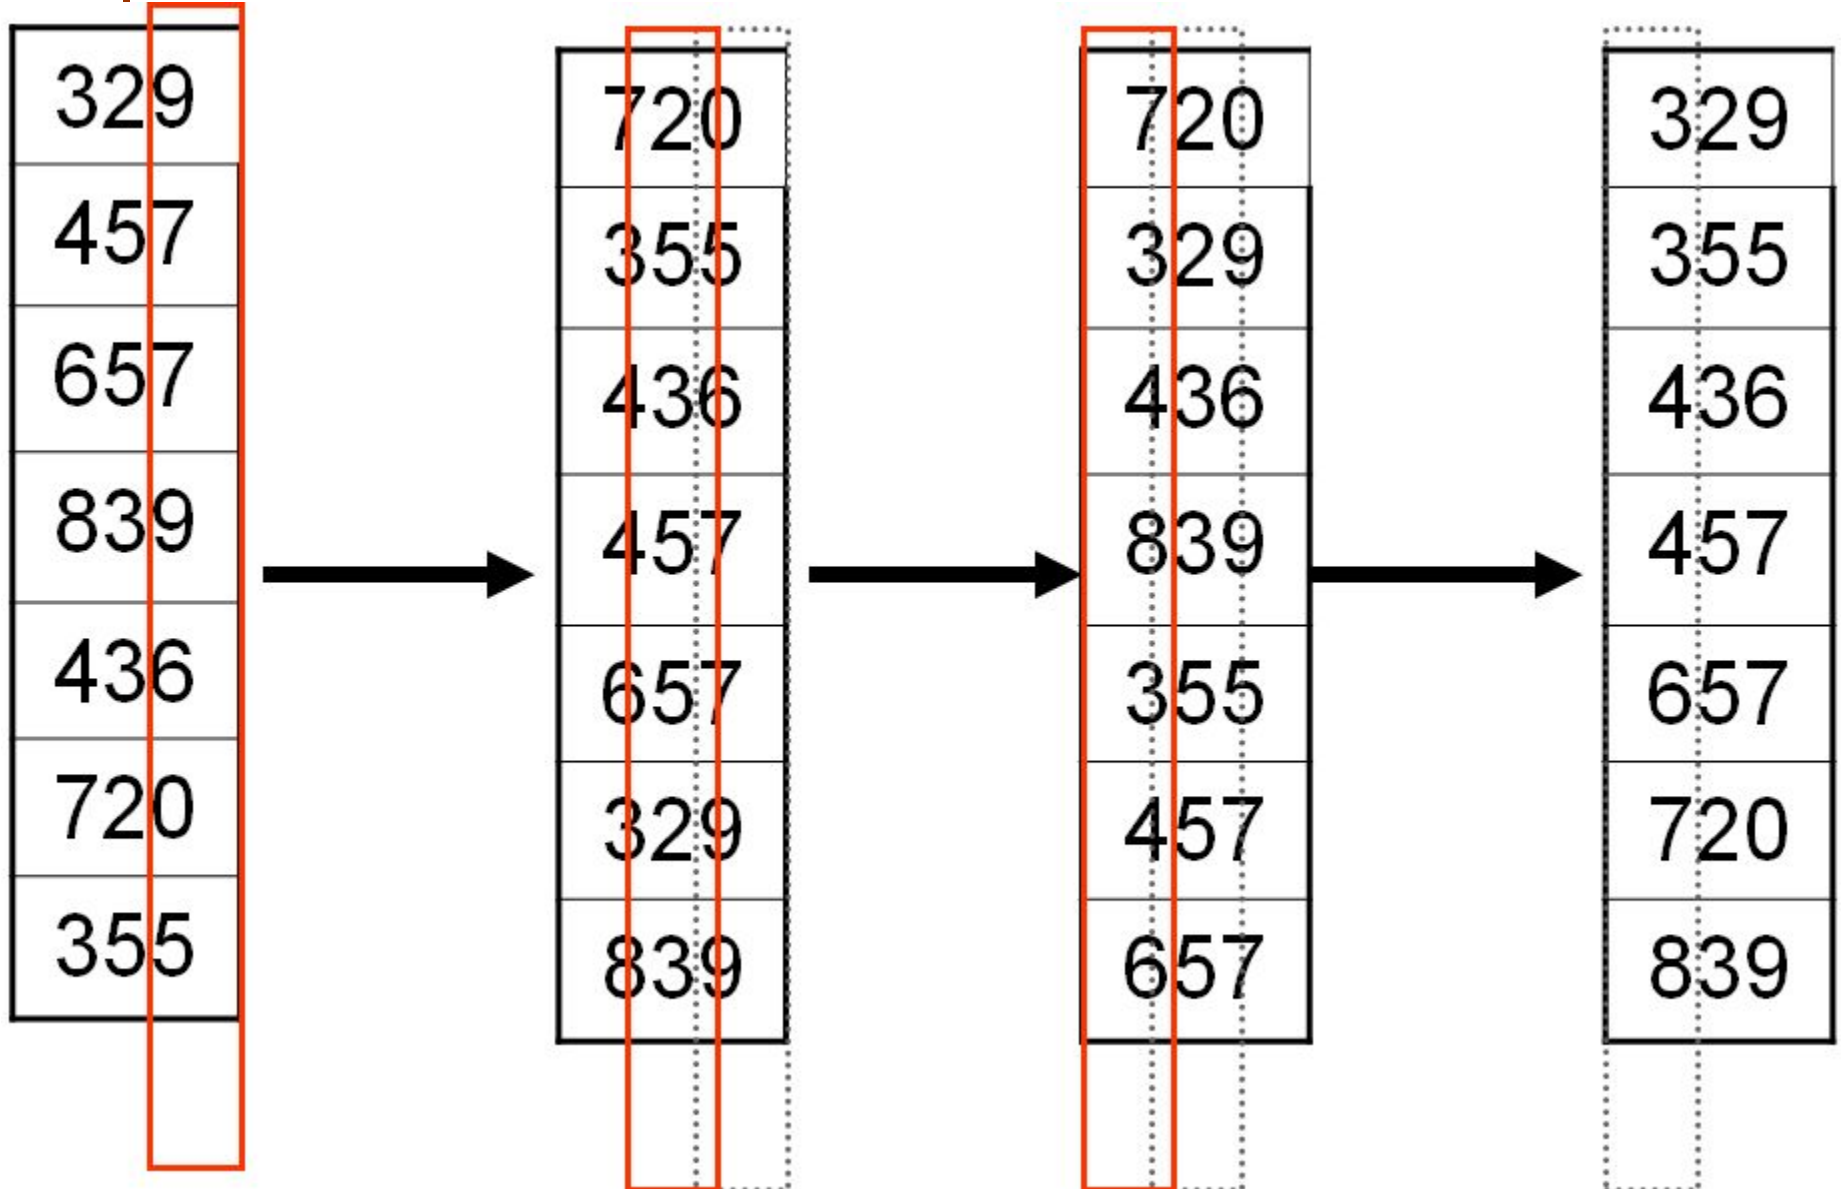
- Intuitively, you might sort on the most significant digit, then the second msd, etc.
- Problem: lots of intermediate piles of cards (read: scratch arrays) to keep track of
- Key idea: sort the least significant digit first

RadixSort(A, d)

for $i=1$ to d

StableSort(A) on digit i

Radix Sort Example





Radix Sort

- What is the running time of radix sort?
 - Each pass over the d digits takes time $O(n+k)$, so total time $O(dn+dk)$
 - When d is constant and $k=O(n)$, takes $O(n)$ time
- Stable, Fast
- Doesn't sort in place (because counting sort is used)



Radix Sort

- ❑ Problem: sort 1 million 64-bit numbers
 - ❑ Treat as four-digit radix 216 numbers
 - ❑ Can sort in just four passes with radix sort!
- ❑ Performs well compared to typical $O(n \lg n)$ comparison sort
 - ❑ Approx $\lg(1,000,000)$ 20 comparisons per number being sorted



Radix Sort

- ❑ Assumption: input has d digits ranging from 0 to k
- ❑ Basic idea:
 - ❑ Sort elements by digit starting with least significant
 - ❑ Use a stable sort (like counting sort) for each stage
- ❑ Pro's:
 - ❑ Fast
 - ❑ Asymptotically fast (i.e., $O(n)$ when d is constant and $k=O(n)$)
 - ❑ Simple to code
 - ❑ A good choice
- ❑ Con's:
 - ❑ Doesn't sort in place
 - ❑ Not a good choice for floating point numbers or arbitrary strings.



Bucket Sort

- Assumption: input - n real numbers from $[0, 1)$
- Basic idea:
 - Create n linked lists (buckets) to divide interval $[0, 1)$ into subintervals of size $1/n$
 - Add each input element to appropriate bucket and sort buckets with insertion sort
- Uniform input distribution $O(1)$ bucket size
 - Therefore the expected total time is $O(n)$

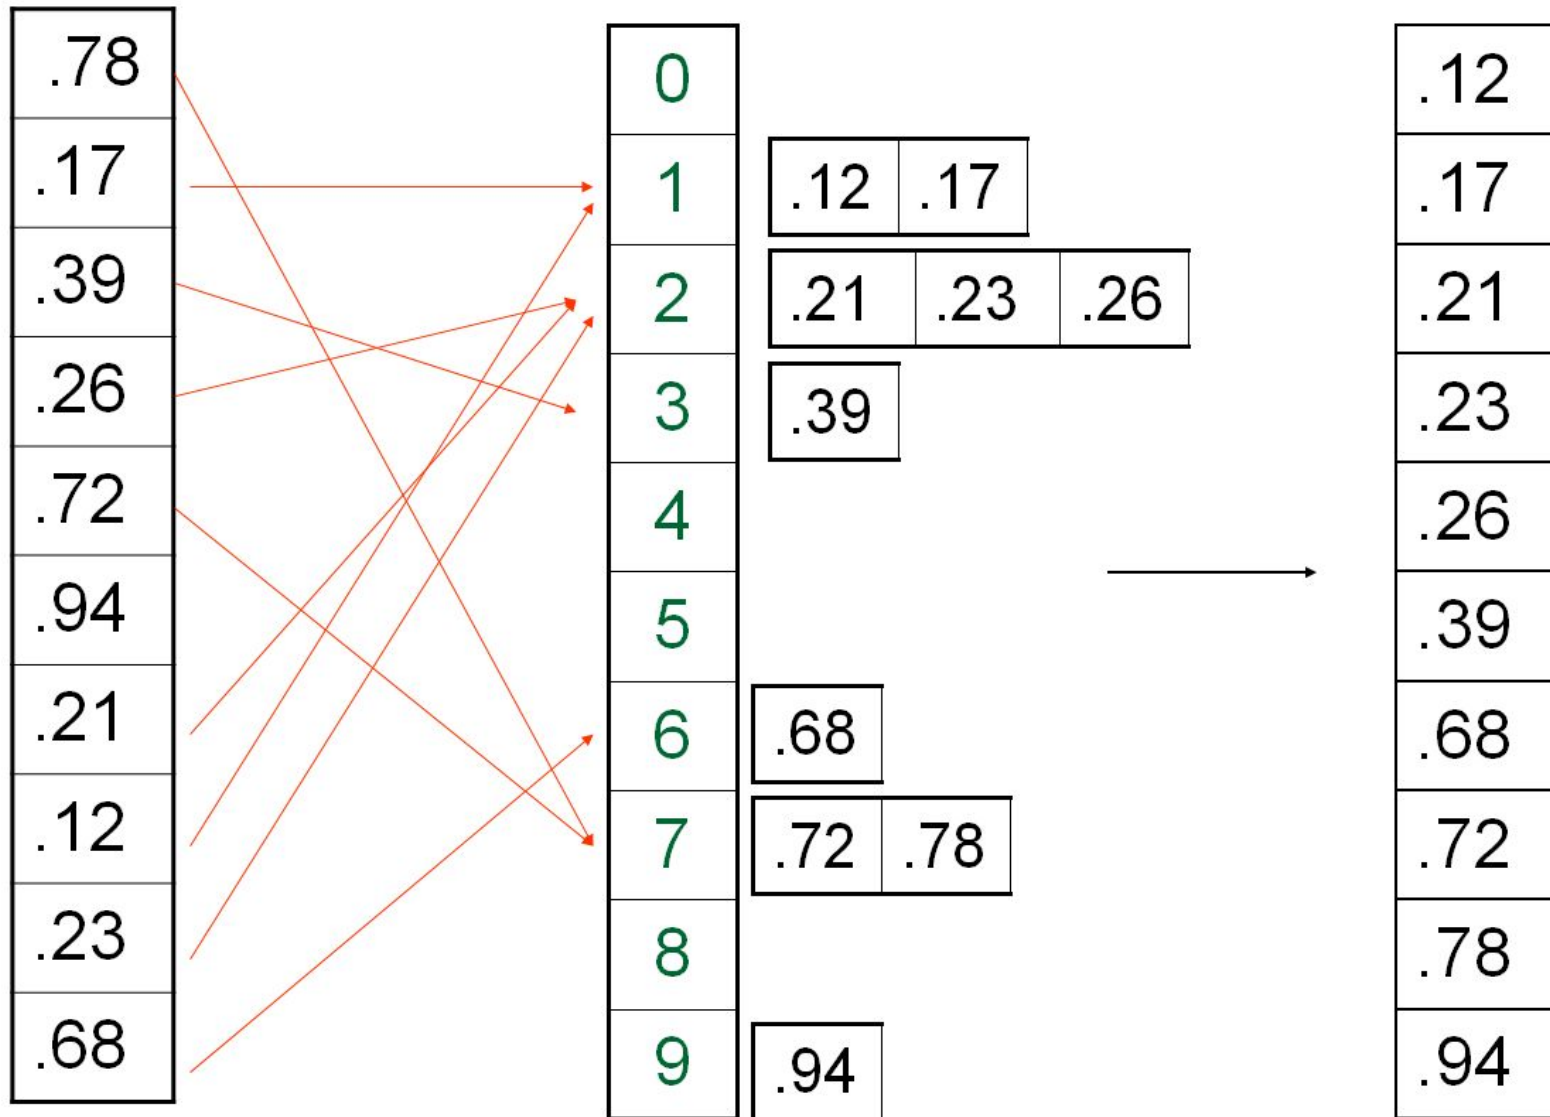


Bucket Sort

Bucket-Sort(A)

1. $n \leftarrow \text{length}(A)$
2. **for** $i \leftarrow 0$ to n \longleftarrow *Distribute elements over buckets*
3. **do** insert $A[i]$ into list $B[\text{floor}(n \cdot A[i])]$
4. **for** $i \leftarrow 0$ to $n - 1$ \longleftarrow *Sort each bucket*
5. **do** Insertion-Sort($B[i]$)
6. Concatenate lists $B[0], B[1], \dots, B[n - 1]$ in order

Bucket Sort Example





Bucket Sort-Running Time

- All lines except line 5 (Insertion-Sort) take $O(n)$ in the worst case.
- In the worst case, $O(n)$ numbers will end up in the same bucket, so in the worst case, it will take $O(n^2)$ time.
- Lemma: Given that the input sequence is drawn uniformly at random from $[0,1)$, the expected size of a bucket is $O(1)$.
- So, in the average case, only a constant number of elements will fall in each bucket, so it will take $O(n)$ (see proof in book).
- Use a different indexing scheme (hashing) to distribute the numbers uniformly.



Bucket Sort Review

- ❑ Assumption: input is uniformly distributed across a range
- ❑ Basic idea:
 - ❑ Partition the range into a fixed number of buckets.
 - ❑ Toss each element into its appropriate bucket.
 - ❑ Sort each bucket.
- ❑ Pro's:
 - ❑ Fast
 - ❑ Asymptotically fast (i.e., $O(n)$ when distribution is uniform)
 - ❑ Simple to code
 - ❑ Good for a rough sort.
- ❑ Con's:
 - ❑ Doesn't sort in place



Summary of Linear Sorting

Non-Comparison Based Sorts

Running Time

	worst-case	average-case	best-case	in place
Counting Sort	$O(n + k)$	$O(n + k)$	$O(n + k)$	no
Radix Sort	$O(d(n + k'))$	$O(d(n + k'))$	$O(d(n + k'))$	no
Bucket Sort		$O(n)$		no

Counting sort assumes input elements are in range $[0, 1, 2, \dots, k]$ and uses array indexing to count the number of occurrences of each value.

Radix sort assumes each integer consists of d digits, and each digit is in range $[1, 2, \dots, k']$.

Bucket sort requires advance knowledge of input distribution (sorts n numbers uniformly distributed in range in $O(n)$ time).